

# **R Bootcamp Part 3**

Dani Navarro  
Amy Perfors

# Today's Plan

1. Loops
2. Branches
3. Functions
4. Programming
5. File system




Loops

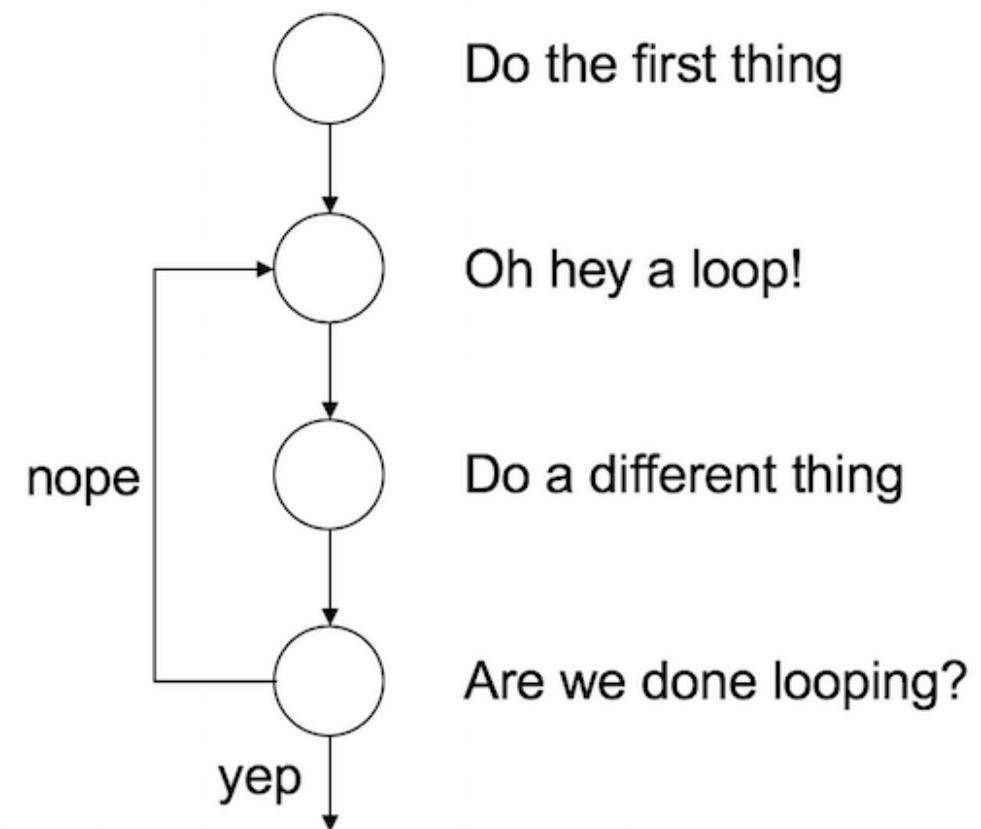
The image features three concentric rectangular frames. The outermost frame is a dark blue line, the middle frame is a medium blue line, and the innermost frame is a light blue line. The word "Loops" is centered within the innermost frame in a black, sans-serif font.

# The purpose of a loop

How R reads a script without loops



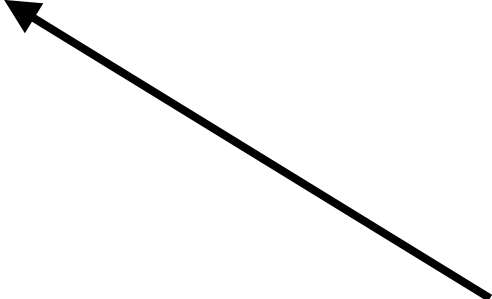
How R reads a script with a loop



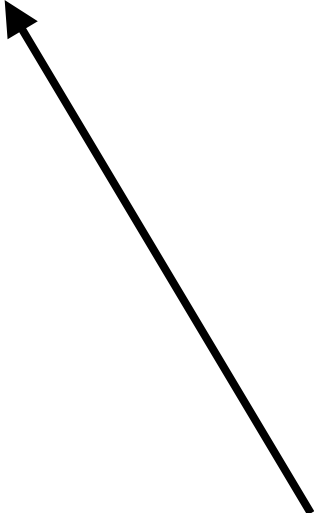
# While loops

```
while ( CONDITION ) {  
    STATEMENT1  
    STATEMENT2  
    ETC  
}
```

Needs to be a logical  
(TRUE or FALSE)

A black arrow points from the text 'Needs to be a logical (TRUE or FALSE)' to the 'CONDITION' part of the while loop syntax.

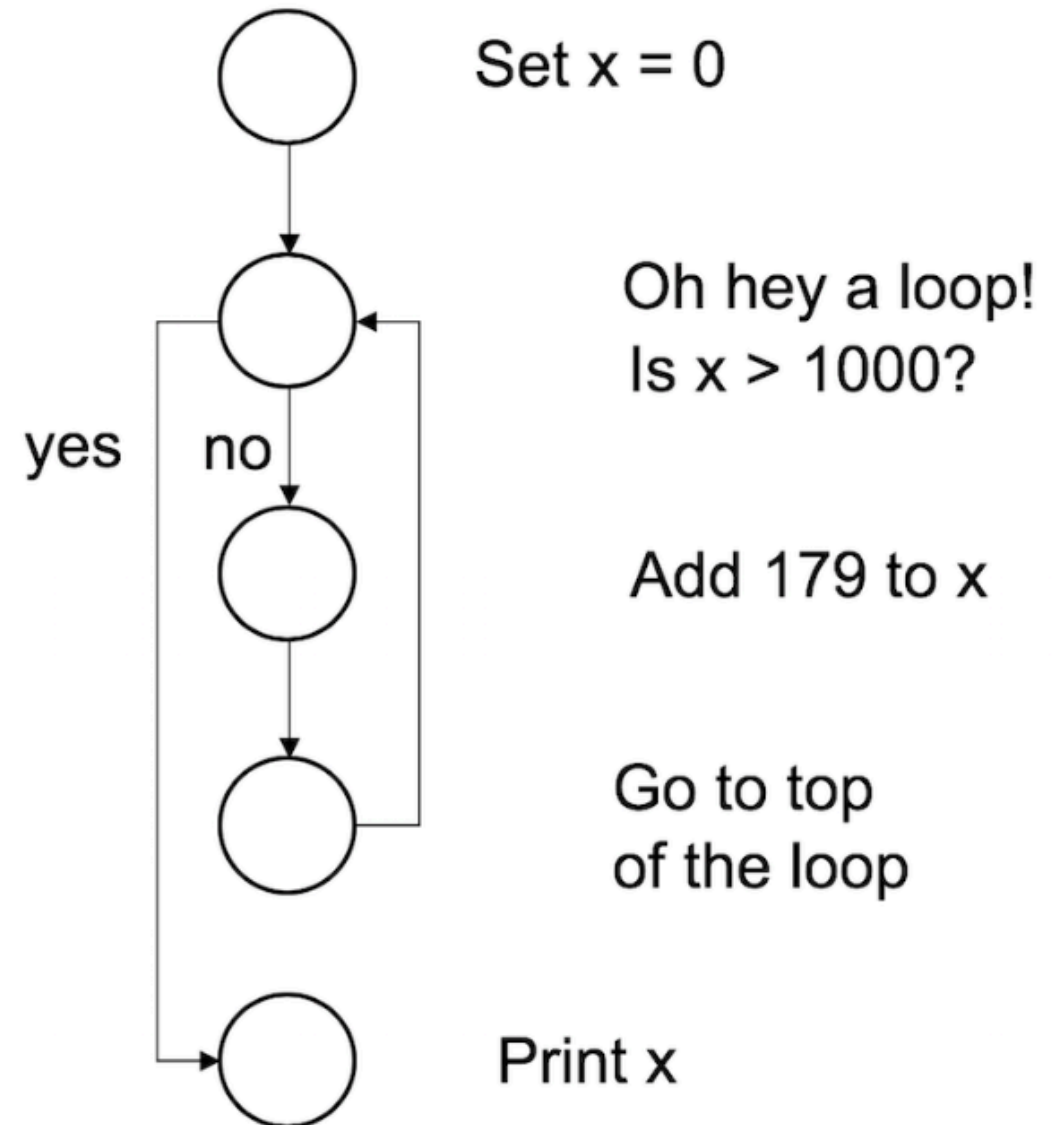
Does all of these  
things as long as the  
condition is TRUE

A black arrow points from the text 'Does all of these things as long as the condition is TRUE' to the body of the while loop (the statements inside the curly braces).

# While loops

```
x <- 0
while (x < 1000) {
  x <- x + 179
}
print(x)
```

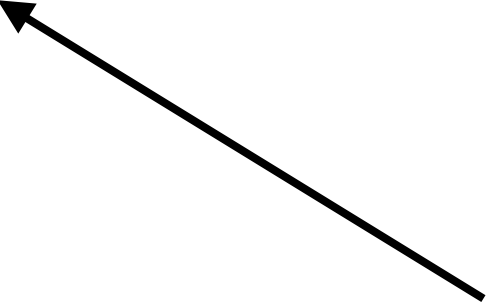
```
## [1] 1074
```



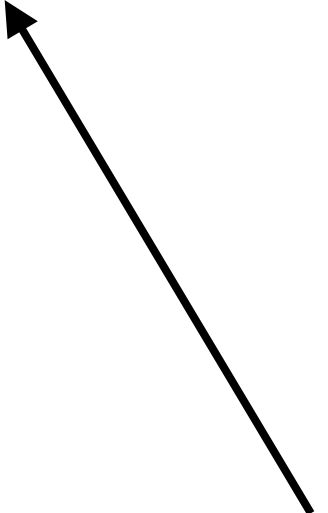
# For loops

```
for ( VAR in VECTOR ) {  
    STATEMENT1  
    STATEMENT2  
    ETC  
}
```

Counts through each  
of the things in the  
vector



Does all of these  
things as long as the  
condition is TRUE



# For loops

```
for ( value in 1:10 ) {  
  answer <- 137*value  
  print(answer)  
}
```

# 137

# 274

# 411

# 548

# 685

# 822

# 959

# 1096

# 1233

# 1370



# Looping over vectors

```
words <- c("farewell", "cruel", "world")
for (thisWord in words) {
  nLetters <- nchar(thisWord)
  blockWord <- toupper(thisWord)
  cat(blockWord, "has", nLetters, "letters\n")
}
```

```
# FAREWELL has 8 letters
# CRUEL has 5 letters
# WORLD has 5 letters
```



Branches

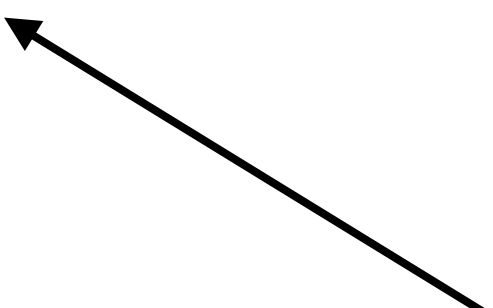
The image features three concentric rectangular frames. The outermost frame is a dark blue line, the middle one is a medium blue line, and the innermost one is a light blue line. The word "Branches" is centered within the innermost frame in a black, sans-serif font.

# Branches

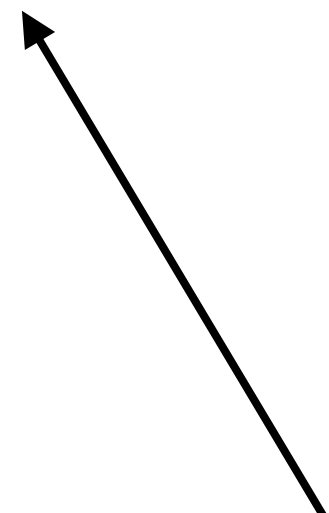
These let you evaluate conditional statements and do different things depending on the outcome

```
If ( CONDITION ) {  
    STATEMENT1  
    STATEMENT2  
    ETC  
}
```

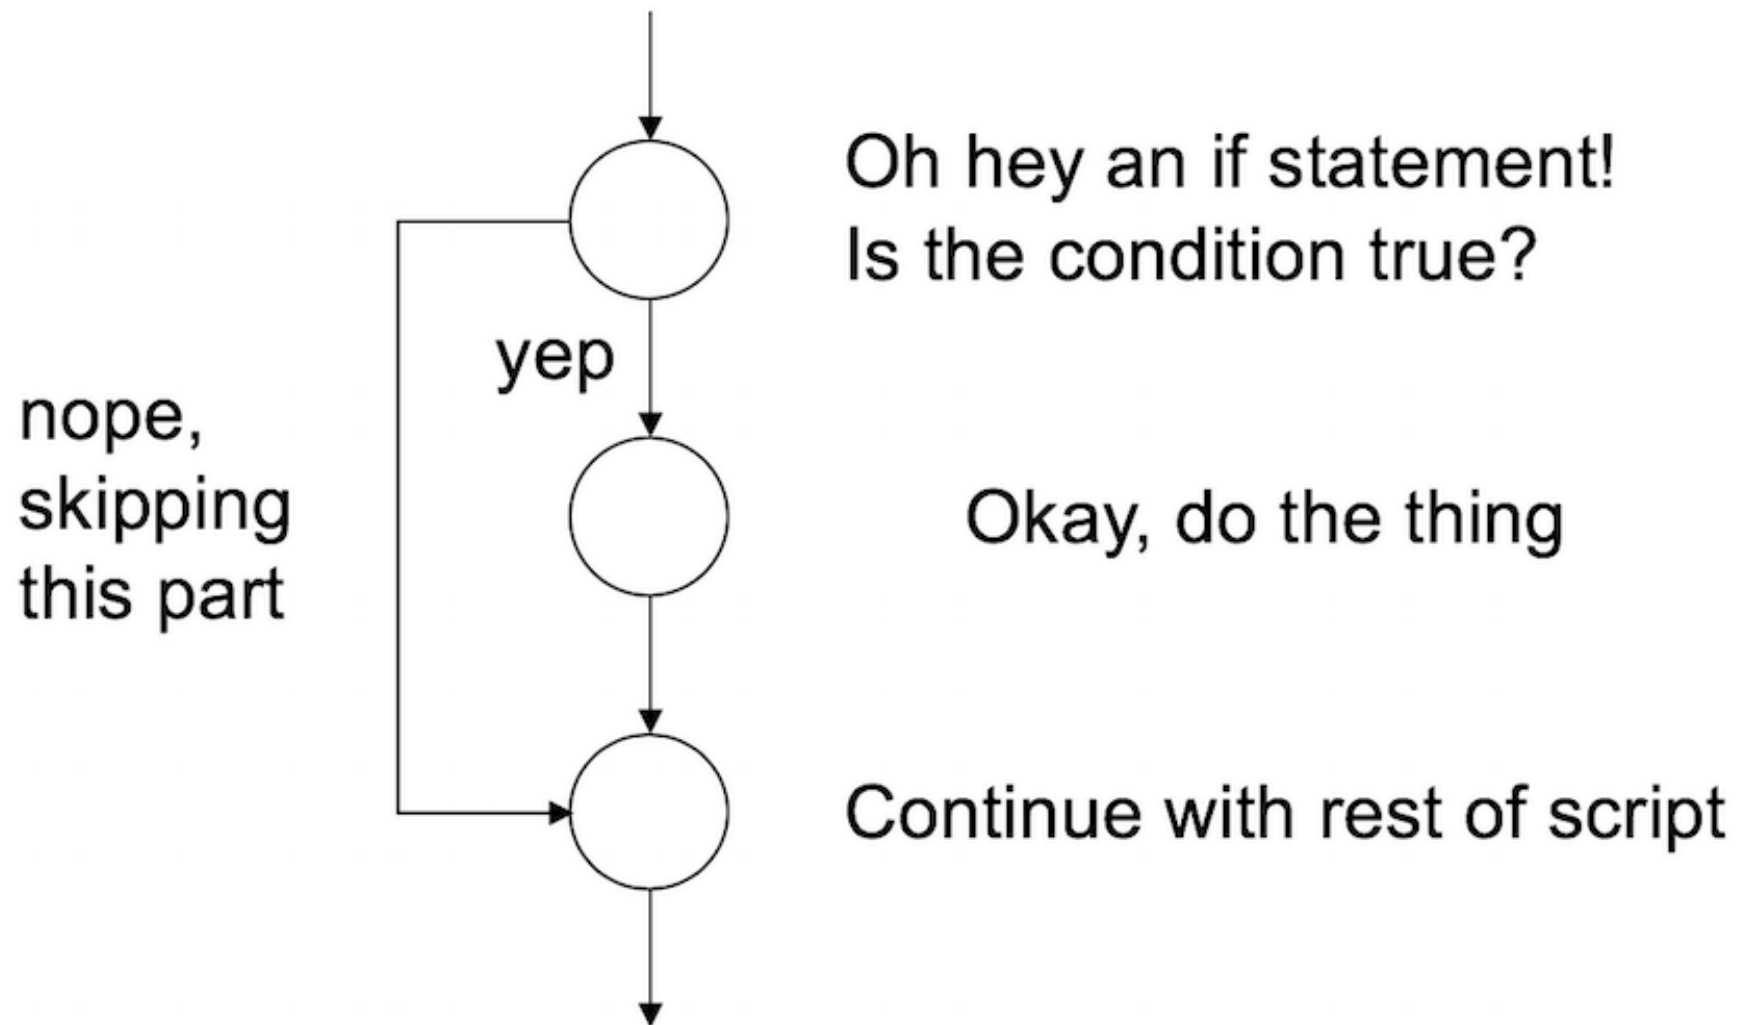
Needs to be a logical  
(TRUE or FALSE)



Does all of these  
things only if the  
condition is TRUE

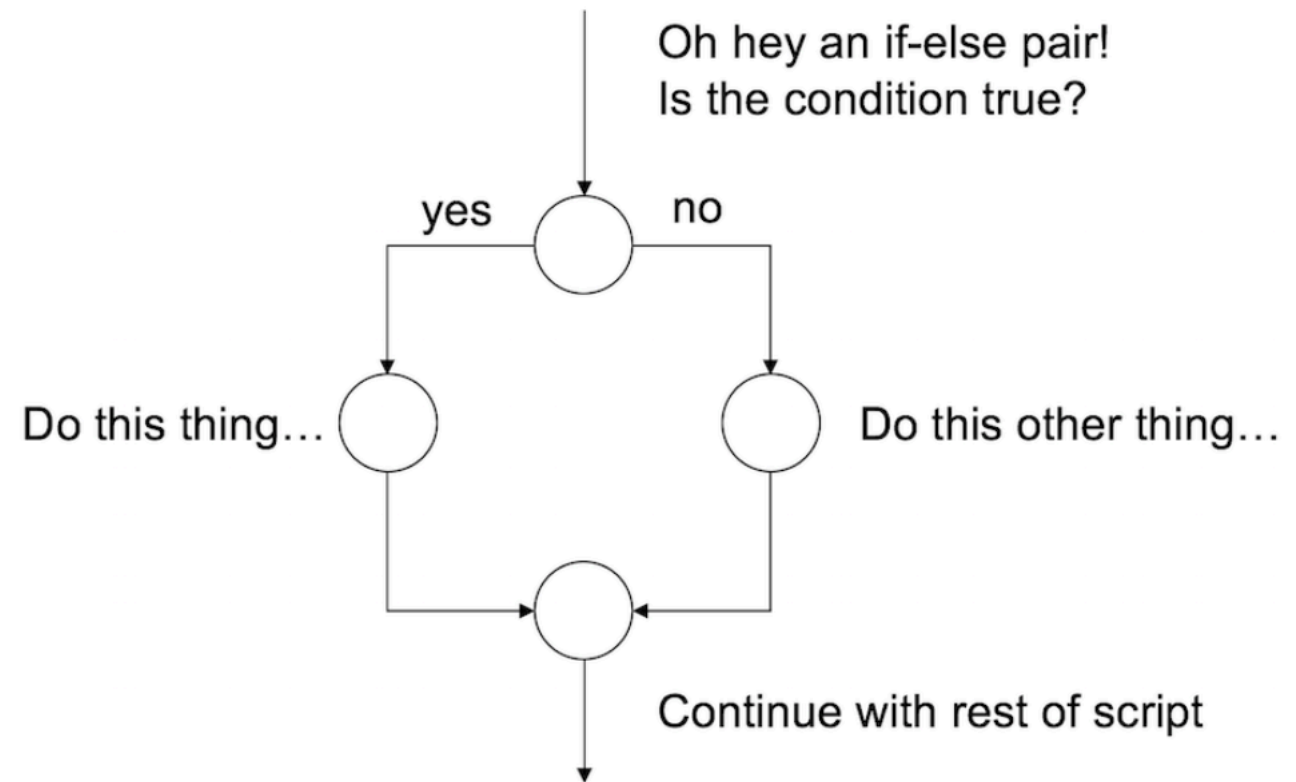


# Branches



# If-Else

```
if ( CONDITION ) {  
    STATEMENT1  
    STATEMENT2  
    ETC  
} else {  
    STATEMENT3  
    STATEMENT4  
}
```



# Example

```
if (today=="Saturday") {  
    print("Yay! Weekend!")  
} else if (today=="Sunday") {  
    print("Uh oh, Monday is coming")  
} else {  
    print("I need coffee.")  
}
```



The image features three concentric rectangular frames. The outermost frame is a dark blue line, the middle frame is a medium blue line, and the innermost frame is a light blue line. The word "Functions" is centered within the innermost frame in a black, sans-serif font.

Functions

# Functions

You can actually create your *own* functions with arguments. Whenever it is called R will execute the statements within it. Creating a function means R creates a temporary environment with it while it's in practice, and only “keeps” the value in the `return()` statement.

```
FNAME <- function (ARG1, ARG2, ARG3, ETC) {  
  STATEMENT1  
  STATEMENT2  
  STATEMENT3  
  ETC  
  return (VALUE)  
}
```



# Functions

Here's an example of a function that will square any number.

```
square <- function(x) {  
  y <- x*x  
  return(y)  
}
```

```
> square(4)  
# 16
```

# Functions

The ... argument lets the user enter as many arguments as they would like, as in the example below.

```
doubleMax <- function(...) {  
  maxVal <- max(...)  
  out <- 2*maxVal  
  return(out)  
}
```



Bringing it all together

# What is all this about?????

Suppose we present a compound stimulus AB, which consists of two things, a tone (A) and a light (B). This compound is presented together with a shock. In associative learning studies, this kind of trial is denoted AB+ to indicate that the outcome (US) was present at the same time as the two stimuli that comprise the CS. According to the Rescorla-Wagner model, the rule for updating the associative strengths  $v_A$  and  $v_B$  between the originally neutral stimuli and the shock is given by:

$$\begin{aligned}v_A &\leftarrow v_A + \alpha_A \beta_U (\lambda_U - v_{AB}) \\v_B &\leftarrow v_B + \alpha_B \beta_U (\lambda_U - v_{AB})\end{aligned}$$

where the associative value  $v_{AB}$  of the compound stimulus AB is just the sum of the values of the two items individually. This is expressed as:


$$v_{AB} = v_A + v_B$$

To understand this rule, note that:

- $\lambda_U$  is a variable that represents the “reward value” (or “punishment value”) of the US itself, and as such represents the maximum possible association strength for the CS.
- $\beta_U$  is a learning rate linked to the US (e.g. how quickly do I learn about shocks?)
- $\alpha_A$  is a learning rate linked to the CS (e.g. how quickly do I learn about tones?)
- $\alpha_B$  is also a learning rate linked to the CS (e.g. how quickly do I learn about lights?)



# Associative learning




In the simplest design (forward conditioning) the CS is presented slightly before the US, so that it can serve as a signal that reward is coming

Unconditioned stimulus (US) –  
something inherently rewarding

Conditioned stimulus (CS) –  
something initially neutral

# Associative learning


(well, Pavlovian anyway)




After some number of presentations, the learner starts to respond to the CS in the same way they would respond to the US

They have a learned association between the CS and the US


## FORWARD CONDITIONING




## SIMULTANEOUS CONDITIONING




## SECOND ORDER CONDITIONING




## TEMPORAL CONDITIONING





## EXTINCTION



## BLOCKING



## INHIBITION




There are many variations on this!

(Long list of empirical effects to account for)

# The Rescorla-Wagner model

$$V_x \leftarrow V_x + \alpha\beta(\lambda - V_{xy})$$



One popular (though flawed & incomplete) account of associative learning is the Rescorla-Wagner model

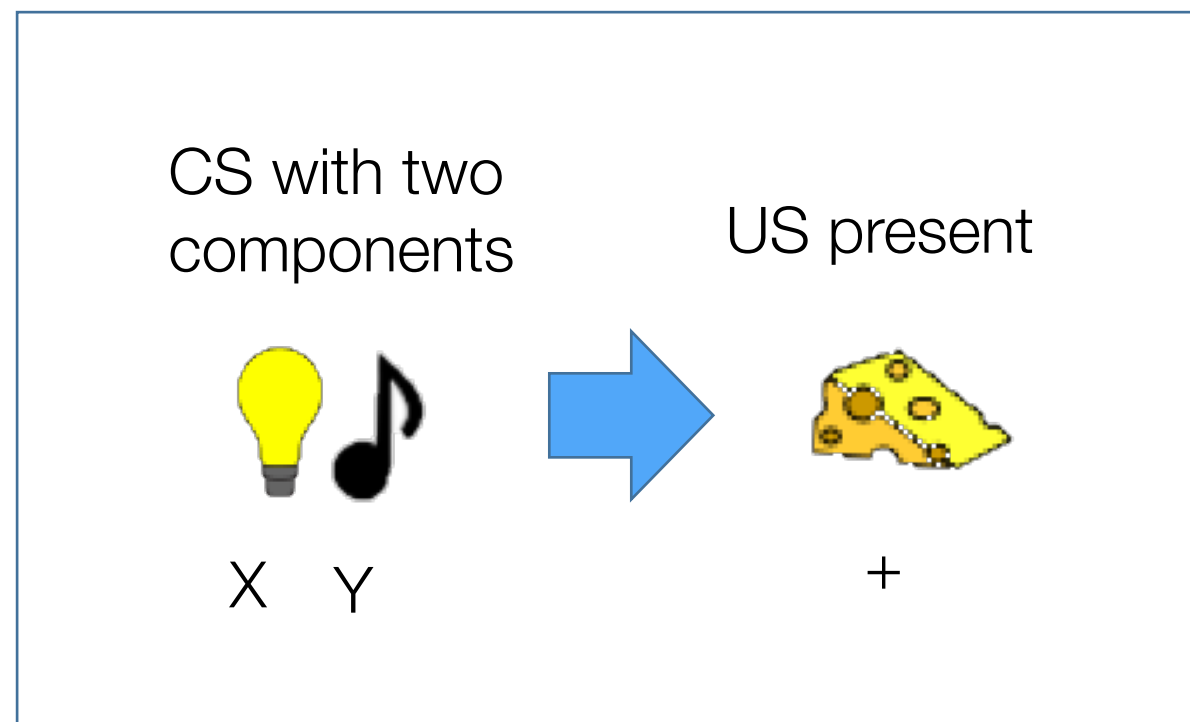


But what does this strange inscription mean?????



# The Rescorla-Wagner model

$$V_x \leftarrow V_x + \alpha\beta(\lambda - V_{xy})$$



An XY+ trial

Consider a design in which there are two features present (X and Y) and the learner needs to predict an outcome that might be present (+) or absent (-)

# The Rescorla-Wagner model

$$V_x \leftarrow V_x + \alpha\beta(\lambda - V_{xy})$$

The old strength of  
association for stimulus X



The new strength of association  
for stimulus X after seeing XY+



The difference between the old and the new.  
By convention “differences” are denoted  
“delta”, so we call this “delta-V”,  $\Delta V$

# The Rescorla-Wagner model

$$\Delta V_x = \alpha \beta (\lambda - V_{xy})$$

This delta-V describes “how much we learn about X from the current trial/event”

The “alpha” and “beta” terms here are parameters describing learning rates.

- alpha depends on the CS
- beta depends on the US



# The Rescorla-Wagner model


$$\Delta V_x = \alpha \beta (\lambda - V_{xy})$$



This difference term here is called  
the “reward prediction error”

# The Rescorla-Wagner model

$$\Delta V_x = \alpha \beta (\lambda - V_{xy})$$



lambda is represents the “intrinsic” value of the outcome (unconditioned stimulus), sometimes referred to as the “reward”, r

# The Rescorla-Wagner model

$$\Delta V_x = \alpha\beta(\lambda - V_{xy})$$

$V_{xy}$  is the “predicted reward”: the amount of reward/punishment that the learner expects to receive upon seeing the compound stimulus XY

In the Rescorla-Wagner model, expectations are additive, which means that:

$$V_{xy} = V_x + V_y$$

(But not all learning models assume additivity)

# Error driven learning!

$$\Delta V_x = \alpha \beta (\lambda - V_{xy})$$

reward

expected reward

the difference between outcomes and expectations is the prediction error, and it is this error that “drives” learning

how much does the learner change their beliefs?

learning is gradual, and depends on a learning rate



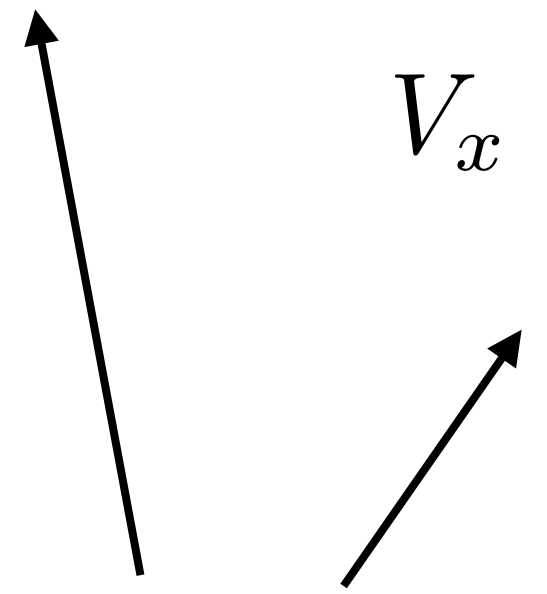
Let's do this!



# Step 1: Skeleton

```
updateRW <- function(value, alpha, beta, lambda) {  
}
```

$$V_x \leftarrow V_x + \alpha\beta(\lambda - V_{xy})$$



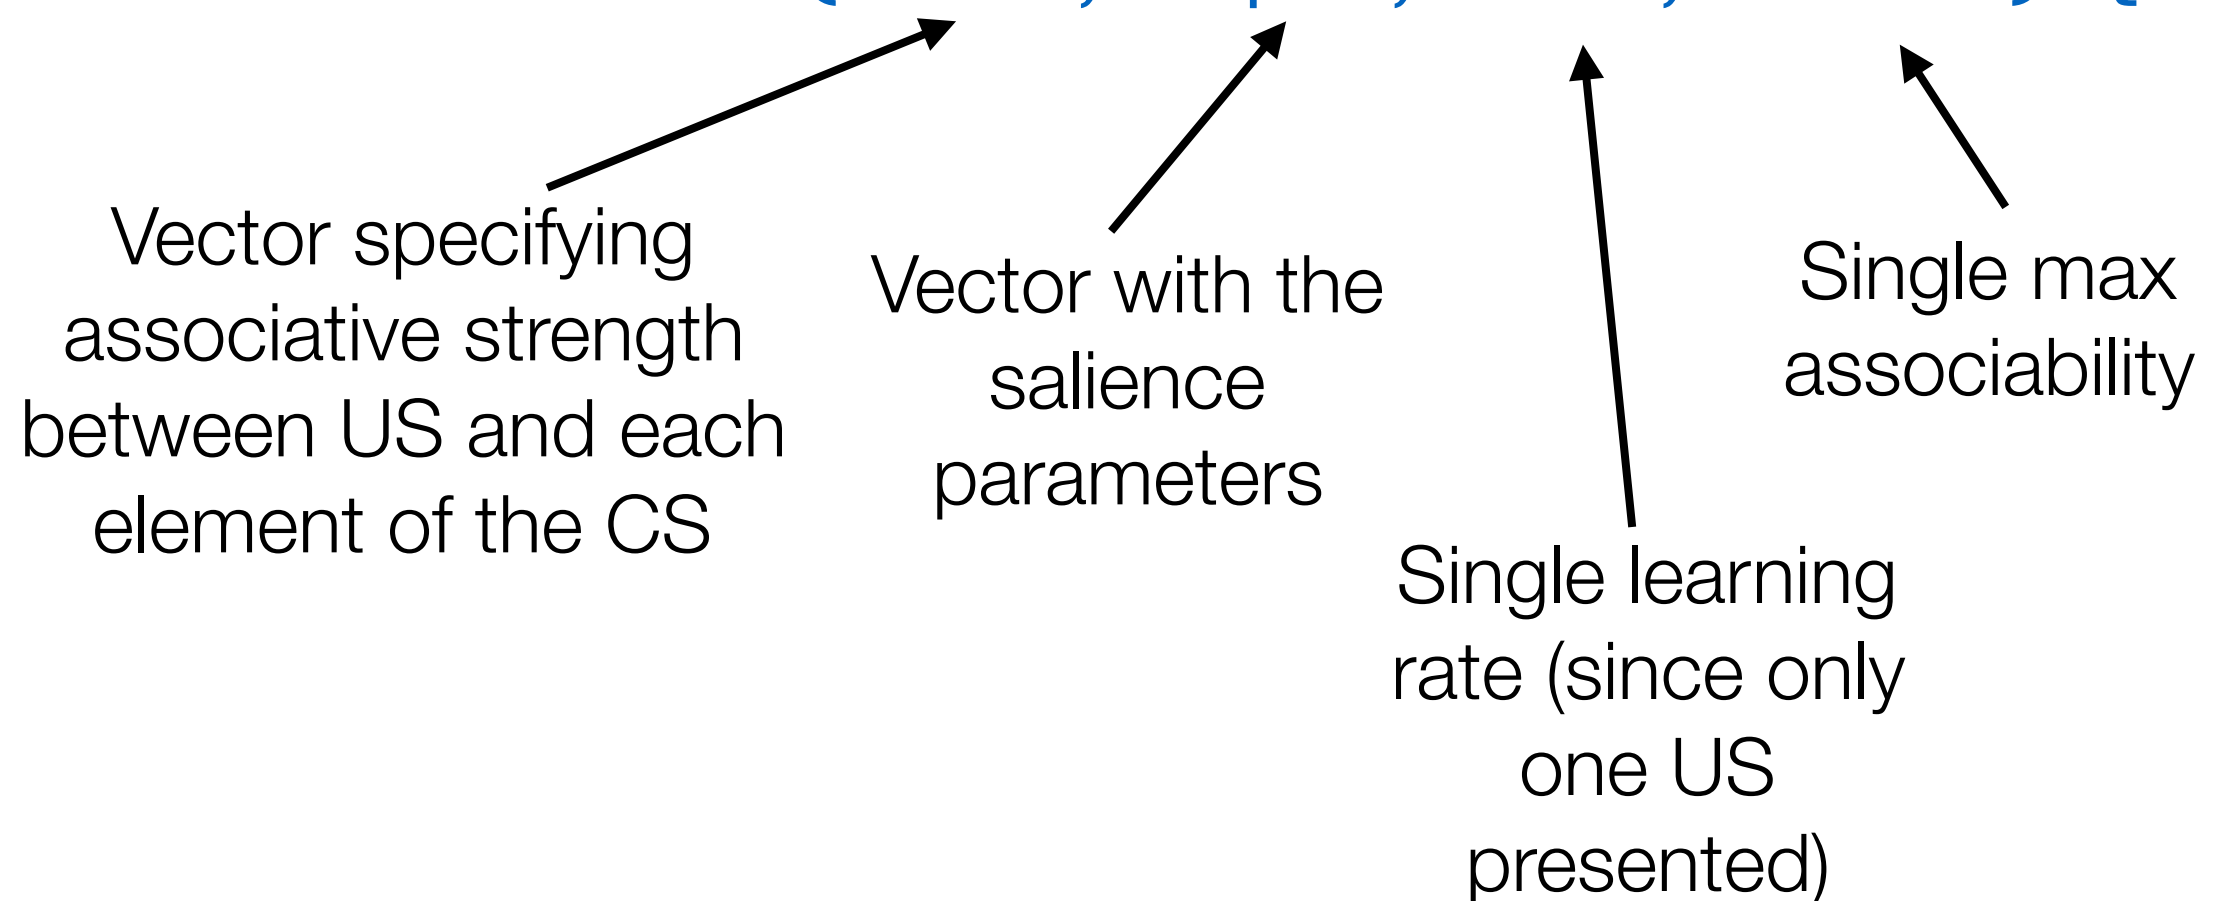
The diagram consists of two black arrows. One arrow originates from the text 'The design of our R function mirrors the structure of the Rescorla-Wagner model that it implements' and points upwards and to the left towards the opening curly brace of the R function definition. The second arrow originates from the same text and points upwards and to the right towards the mathematical equation  $V_x \leftarrow V_x + \alpha\beta(\lambda - V_{xy})$ , which is positioned between the two arrows.

The design of our R function  
mirrors the structure of the  
Rescorla-Wagner model that it  
implements

# Step 1: Skeleton

```
updateRW <- function(value, alpha, beta, lambda) {  
}
```

Vector specifying  
associative strength  
between US and each  
element of the CS



Vector with the  
salience  
parameters

Single learning  
rate (since only  
one US  
presented)

Single max  
associability

Reminder:

$$V_x \leftarrow V_x + \alpha\beta(\lambda - V_{xy})$$

## Step 2: Make a plan

```
updateRW <- function(value, alpha, beta, lambda) {  
  
  # compute the value of the compound stimulus  
  # compute the prediction error  
  # compute the change in strength  
  # update the association value  
  # return the new value  
  
}
```

Reminder:

$$V_x \leftarrow V_x + \alpha\beta(\lambda - V_{xy})$$

## Step 3: Put in the details

```
updateRW <- function(value, alpha, beta, lambda) {  
  # compute the value of the compound stimulus  
  valueCompound <- sum(value)  
  
  # compute the prediction error  
  predictionError <- lambda - valueCompound  
  
  # compute the change in strength  
  valueChange <- alpha * beta * predictionError  
  
  # update the association value  
  value <- value + valueChange  
  
  # return the new value  
  return(value)  
}
```

Reminder:


$$V_x \leftarrow V_x + \alpha\beta(\lambda - V_{xy})$$

# Step 4: Model predictions

1. Conditioning
2. Extinction
3. Blocking

# Conditioning

```
nTrials <- 20  
strength <- numeric(nTrials)  
  
for (trial in 2:nTrials) {  
  strength[trial] <- updateRW(strength[trial-1])  
}
```



# Extinction


```
nTrials <- 50  
strength <- numeric(nTrials)  
lambda <- 0.3
```

```
for (trial in 2:nTrials) {
```

```
  # remove the shock after trial 25  
  if (trial>25) {  
    lambda <- 0  
  }
```

```
  # update associative strength on each trial  
  strength[trial] <- updateRW(value=strength[trial-1],  
                              lambda=lambda)
```

```
}
```



# Blocking

```
# total number of trials across  
# both phases of the task  
n_trials <- 50  
  
# vectors of zeros  
strength_A <- rep(0,n_trials)  
strength_B <- rep(0,n_trials)
```




# Blocking

```
# total number of trials across  
# both phases of the task  
n_trials <- 50  
  
# vectors of zeros  
strength_A <- rep(0,n_trials)  
strength_B <- rep(0,n_trials)  
  
# learning rate for the CS at the  
# start of the experiment is .3 for  
# A and 0 for B (b/c it's absent)  
alpha <- c(.3, 0)
```

# Blocking

```
for(trial in 2:n_trials) {  
  
  # after trial 15, both stimuli are present  
  if(trial > 15) alpha <- c(.3, .3)  
  
  # vector of current associative strengths  
  v_old <- c(strength_A[trial-1], strength_B[trial-1])  
  
  # vector of new associative strengths  
  v_new <- update_RW(  
    value = v_old,  
    alpha = alpha  
  )  
  
  # record the new strengths  
  strength_A[trial] <- v_new[1]  
  strength_B[trial] <- v_new[2]  
}
```

# Blocking



Strong association to A is formed early and maintained

There is learning to B, but greatly reduced and it asymptotes at a low level

# Intro to R cheat sheet

## 1 Saving and importing

- Save as .RData, using menu or `save.image()`
- Can load .csv, using menu or `read.csv()`

## 12 Scripts let you run and save series of commands



13

`help(functionName)`  
e.g. `help(print)`

Files Plots Packages Help Viewer

R: Print Values Find in Topic

print {base}

### Print Values

#### Description

print prints its argument and returns it in a function which means that new printing method

#### Usage

```
print(x, ...)
```

#### Arguments

x	an object used to select
...	further arguments passed
quote	logical, indicating whether quotes.
max.levels	integer, indicating how extra "Levels" line will max.levels such that
width	only used when max.levels